

# Lab 1

August 11, 2016

**Part 1.** We will work through some details on logistic regression. We first set some notation. Let  $x_i \in \mathbb{R}^p, y_i \in \{0, 1\}$  for  $i = 1, \dots, n$ . Let  $X \in \mathbb{R}^{n \times p}$  denote the matrix with  $x_i$  as its  $i$ th row. Recall that the negative log-likelihood  $\ell(\beta)$  can be written as follows:

$$\ell(\beta) = \sum_{i=1}^n [-y_i x_i^T \beta + \log(1 + \exp(x_i^T \beta))].$$

1. Write the gradient and Hessian of  $\ell(\beta)$ .
2. What is the computational complexity for calculating the gradient and Hessian of  $\ell(\beta)$ ?
3. Under what condition is  $\ell(\beta)$  strictly convex?
4. Prove that  $\ell(\beta)$  is  $L$ -Lipschitz differentiable with  $L = \frac{1}{4} \|X\|_{\text{op}}^2$ .
5. Suppose that there is a vector  $w \in \mathbb{R}^p$  such that  $x_i^T w > 0$  if  $y_i = 1$  and  $x_i^T w < 0$  if  $y_i = 0$ . Prove that  $\ell(\beta)$  does not have a global minimum. In other words, when the classification problem is completely separable, there is no maximum likelihood estimator.

## Part 2. Gradient Descent

You will now implement gradient descent in your R package. Your function will include using both a fixed step-size as well as one chosen by backtracking.

Please complete the following steps.

**Step 1:** Write a function “gradient\_step.”

```
#' Gradient Step
#'
#' @param gradf handle to function that returns gradient of objective function
#' @param x current parameter estimate
#' @param t step-size
gradient_step <- function(gradf, t, x) {
}

```

Your function should return  $x^+ = x - t\nabla f(x)$ .

**Step 2:** Write a function “gradient\_descent\_fixed.” The function should terminate either when a maximum number of iterations has been taken or if the relative change in the objective function has fallen below a tolerance

$$\frac{|f(x_k) - f(x_{k-1})|}{|f(x_{k-1})| + 1} < \text{tolerance}$$

```
#' Gradient Descent (Fixed Step-Size)
#'
#' @param fx handle to function that returns objective function values
#' @param gradf handle to function that returns gradient of objective function
#' @param x0 initial parameter estimate
#' @param t step-size
#' @param max_iter maximum number of iterations
#' @param tol convergence tolerance
gradient_descent_fixed <- function(fx, gradf, t, x0, max_iter=1e2, tol=1e-3) {
}

```

Your function should return

- The final iterate value
- The objective function values
- The 2-norm of the gradient values
- The relative change in the function values
- The relative change in the iterate values

**Step 3:** Write a function “backtrack.”

```
## Backtracking
##
## @param fx handle to function that returns objective function values
## @param x current parameter estimate
## @param t current step-size
## @param df the value of the gradient of objective function evaluated at the current x
## @param alpha the backtracking parameter
## @param beta the decremting multiplier
backtrack <- function(fx, t, x, df, alpha=0.5, beta=0.9) {
}

```

Your function should return the selected step-size.

**Step 4:** Write a function “gradient\_descent\_backtrack” that performs gradient descent using backtracking.

```
## Gradient Descent (Backtracking Step-Size)
##
## @param fx handle to function that returns objective function values
## @param gradf handle to function that returns gradient of objective function
## @param x0 initial parameter estimate
## @param max_iter maximum number of iterations
## @param tol convergence tolerance
gradient_descent_backtrack <- function(fx, gradf, x0, max_iter=1e2, tol=1e-3) {
}

```

Your function should return

- The final iterate value
- The objective function values
- The 2-norm of the gradient values
- The relative change in the function values
- The relative change in the iterate values

**Step 5:** Write a function “gradient\_descent” that is a wrapper function for “gradient\_descent\_fixed” and “gradient\_descent\_backtrack.” The default should be to use the backtracking.

```
#' Gradient Descent
#'
#' @param fx handle to function that returns objective function values
#' @param gradf handle to function that returns gradient of objective function
#' @param x0 initial parameter estimate
#' @param t step-size
#' @param max_iter maximum number of iterations
#' @param tol convergence tolerance
gradient_descent <- function(fx, gradf, x0, t=NULL, max_iter=1e2, tol=1e-3) {
}

```

Your function should return

- The final iterate value
- The objective function values
- The 2-norm of the gradient values
- The relative change in the function values
- The relative change in the iterate values

**Step 6:** Write functions ‘fx\_logistic’ and ‘gradf\_logistic’ to perform ridge logistic regression

```
#' Objective Function for Logistic Regression
#'
#' @param y binary response
#' @param X design matrix
#' @param beta regression coefficient vector
#' @param lambda regularization parameter
fx_logistic <- function(y, X, beta, lambda=0) {
}

#' Gradient for Logistic Regression
#'
#' @param y binary response
#' @param X design matrix
#' @param beta regression coefficient vector
#' @param lambda regularization parameter
gradf_logistic <- function(y, X, beta, lambda=0) {
}


```

**Step 7:** Perform logistic regression (with  $\lambda = 0$ ) on the following data example  $(y, X)$  using the fixed step-size. Use your answers to Part 1 to choose an appropriate fixed step-size. Plot the difference  $\ell(\beta_k) - \ell(\beta_{10000})$  versus the iteration  $k$ . Comment on the shape of the plot given what you know about the iteration complexity of gradient descent with a fixed step size.


```
set.seed(12345)
n <- 100
p <- 2

```


```
X <- matrix(rnorm(n*p),n,p)
beta0 <- matrix(rnorm(p),p,1)
y <- (runif(n) <= plogis(X%*%beta0)) + 0
```



**Step 8:** Perform logistic regression (with  $\lambda = 0$ ) on the simulated data above using backtracking. Plot the difference  $\ell(\beta_k) - \ell(\beta_{10000})$  versus the iteration  $k$ . Comment on the shape of the plot given what you know about the iteration complexity of gradient descent with backtracking.



**Step 9:** Perform logistic regression (with  $\lambda = 10$ ) on the simulated data above using the fixed step-size. Plot the difference  $\ell_\lambda(\beta_k) - \ell_\lambda(\beta_{10000})$  versus the iteration  $k$ . Comment on the shape of the plot given what you know about the iteration complexity of gradient descent applied to strongly convex functions.



**Part 3.** If you have time, try out higher order methods. Install the R package lbfgs. How does the run time compare with gradient descent? You might find the `'system.time()'` command useful. Try implementing Newton's method as discussed in lecture. How does the run time compare with gradient descent? BFGS?